

CHAPTER 23	526
Special Algorithms for Protocols	526
MULTIPLE-KEY PUBLIC-KEY	526
SECRET-SHARING ALGORITHMS	527
LaGrange Interpolating Polynomial.....	527
Vector Scheme	528
Asmuth-Bloom	528
Karnin-Greene-Hellman	529
Advanced Threshold Schemes	529
Sharing a Secret with Cheaters	530
SUBLIMINAL CHANNEL.....	530
Ong-Schnorr-Shamir	530
ElGamal	531
ESZGN	532
DSA	533
Foiling the DSA Subliminal Channel	535
Other Schemes	535
UNDENIABLE DIGITAL SIGNATURES	535
Concoertible Undeniable Signatures	537
DESIGNATED CONFIRMER SIGNATU ...	538
COMPUTING WITH ENCRYPTED	539
The Discrete Logarithm Problem	539
FAIR COIN FLIPS	540
Coin Flipping Using Square Roots	540
Coin Flipping Using Exponentiation	541
Coin Flipping Using Blum Integers	542
ONE-WAY ACCUMULATORS	542
ALL-OR-NOTHING DISCLOSURE OF	542
FAIR AND FAILSAFE CRYPTOSYSTE	545
Fair Diffie-Hellman	545
Failsafe Diffie-Hellman	546
ZERO-KNOWLEDGE PROOFS OF.....	547
Zero-Knowledge Proof a Discrete	547
Zero-Knowledge Proof of the Ability to.....	547
Zero-Knowledge Proof that n Is a Blwm.....	548
BLIND SIGNATURES	548

OBLIVIOUS TRANSFER	549
SECURE MULTIPARTY COMPUTATIO...	550
Example of the Protocol	551
PROBABILISTIC ENCRYPTION.....	551
QUANTUM CRYPTOGRAPHY	553

CHAPTER 23

Special Algorithms for Protocols

23.1 MULTIPLE-KEY PUBLIC-KEY CRYPTOGRAPHY

This is a generalization of RSA (see Section 19.3) [217,212]. The modulus, n , is the product of two primes, p and q . However, instead of choosing e and d such that $ed \equiv 1 \pmod{(p-1)(q-1)}$, choose t keys, K_i , such that

$$K_1 * K_2 * \dots * K_t \equiv 1 \pmod{(p-1)(q-1)}$$

Since

$$M^{K_1 * K_2 * \dots * K_t} = M$$

this is a multiple-key scheme as described in Section 3.5.

If, for example, there are five keys, a message encrypted with K_3 and K_5 can be decrypted with K_1 , K_2 , and K_4 :

$$C = M^{K_3 * K_5} \pmod{n}$$

$$M = C^{K_1 * K_2 * K_4} \pmod{n}$$

One use for this is multisignatures. Imagine a situation where both Alice and Bob have to sign a document for it to be valid. Use three keys: K_1 , K_2 , and K_3 . The first two are issued one each to Alice and Bob, and the third is made public.

- (1) First Alice signs M and sends it to Bob.

$$M' = M^{K_1} \pmod{n}$$

- (2) Bob can recover M from M' .

$$M = M'^{K_2 * K_3} \pmod{n}$$

- (3) He can also add his signature.

$$M'' = M'^{K_2} \pmod{n}$$

- (4) Anyone can verify the signature with K_3 , the public key.

$$M = M''^{K_3} \bmod n$$

Note that a trusted party is needed to set this system up and distribute the keys to Alice and Bob. Another scheme with the same problem is [484]. Yet a third scheme is [695,830,700], but the effort in verification is proportional to the number of signers. Newer schemes [220,1200] based on zero-knowledge identification schemes solve both shortcomings of the previous systems.

23.2 SECRET-SHARING ALGORITHMS

Back in Section 3.7 I discussed the idea behind secret-sharing schemes. The four different algorithms that follow are all particular cases of a general theoretical framework [883].

LaGrange Interpolating Polynomial Scheme

Adi Shamir uses polynomial equations in a finite field to construct a threshold scheme [1414]. Choose a prime, p , which is both larger than the number of possible shadows and larger than the largest possible secret. To share a secret, generate an arbitrary polynomial of degree $m - 1$. For example, if you want to create a $(3,n)$ -threshold scheme (three shadows are necessary to reconstruct M), generate a quadratic polynomial

$$(ax^2 + bx + M) \bmod p$$

where p is a random prime larger than any of the coefficients. The coefficients a and b are chosen randomly; they are kept secret and discarded after the shadows are handed out. M is the message. The prime must be made public.

The shadows are obtained by evaluating the polynomial at n different points:

$$k_i = F(x_i)$$

In other words, the first shadow could be the polynomial evaluated at $x = 1$, the second shadow could be the polynomial evaluated at $x = 2$, and so forth.

Since the quadratic polynomial has three unknown coefficients, a , b , and M , any three shadows can be used to create three equations. Two shadows cannot. One shadow cannot. Four or five shadows are redundant.

For example, let M be 11. To construct a $(3, 5)$ -threshold scheme, where any three of five people can reconstruct M , first generate a quadratic equation (7 and 8 were chosen randomly):

$$F(x) = (7x^2 + 8x + 11) \bmod 13$$

The five shadows are:

$$k_1 = F(1) = 7 + 8 + 11 \equiv 0 \pmod{13}$$

$$k_2 = F(2) = 28 + 16 + 11 \equiv 3 \pmod{13}$$

$$k_3 = F(3) = 63 + 24 + 11 \equiv 7 \pmod{13}$$

$$k_4 = F(4) = 112 + 32 + 11 \equiv 12 \pmod{13}$$

$$k_5 = F(5) = 175 + 40 + 11 \equiv 5 \pmod{13}$$

To reconstruct M from three of the shadows, for example k_2 , k_3 , and k_5 , solve the set of linear equations:

$$a \cdot 2^2 + b \cdot 2 + M \equiv 3 \pmod{13}$$

$$a \cdot 3^2 + b \cdot 3 + M \equiv 7 \pmod{13}$$

$$a \cdot 5^2 + b \cdot 5 + M \equiv 5 \pmod{13}$$

The solution will be $a = 7$, $b = 8$, and $M = 11$. So M is recovered.

This sharing scheme can be easily implemented for larger numbers. If you want to divide the message into 30 equal parts such that any six can get together and reproduce the message, give each of the 30 people the evaluation of a polynomial of degree 5.

$$F(x) = (ax^6 + bx^5 + cx^4 + dx^3 + ex^2 + fx + M) \bmod p$$

Six people can solve for the six unknowns (including M); five people cannot learn anything about M .

The most mind-boggling aspect of secret sharing is that if the coefficients are picked randomly, five people with infinite computing power can't learn anything more than the length of the message (which each of them knows anyway). This is as secure as a one-time pad; an attempt at exhaustive search (that is, trying all possible sixth shadows) will reveal that any conceivable message could be the secret. This is true for all the secret-sharing schemes presented here.

Vector Scheme

George Blakley invented a scheme using points in space [182]. The message is defined as a point in m -dimensional space. Each shadow is the equation of an $(m - 1)$ -dimensional hyperplane that includes the point. The intersection of any m of the hyperplanes exactly determines the point.

For example, if three shadows are required to reconstruct the message, then it is a point in three-dimensional space. Each shadow is a different plane. With one shadow, you know the point is somewhere on the plane. With two shadows, you know the point is somewhere on the line formed where the two planes intersect. With three shadows, you can determine the point exactly: the intersection of the three planes.

Asmuth-Bloom

This scheme uses prime numbers [65]. For an (m, n) -threshold scheme, choose a large prime, p , greater than M . Then choose n numbers less than p , d_1, d_2, \dots, d_n , such that:

1. The d values are in increasing order; $d_i < d_{i+1}$
2. Each d_i is relatively prime to every other d_i
3. $d_1 \cdot d_2 \cdot \dots \cdot d_m > p \cdot d_{n-m+2} \cdot d_{n-m+3} \cdot \dots \cdot d_n$

To distribute the shadows, first choose a random value r and compute

$$M' = M + rp$$

The shadows, k_i , are

$$k_i = M' \bmod d_i$$

Any m shadows can get together and reconstruct M using the Chinese remainder theorem, but any $m - 1$ cannot. See [65] for details.

Karnin-Greene-Hellman

This scheme uses matrix multiplication [818]. Choose $n + 1$ m -dimensional vectors, V_0, V_1, \dots, V_n , such that any possible $m * m$ matrix formed out of those vectors has rank m . The vector U is a row vector of dimension $m + 1$.

M is the matrix product $U \cdot V_0$. The shadows are the products $U \cdot V_i$, where i is a number from 1 to n .

Any m shadows can be used to solve the $m * m$ system of linear equations, where the unknowns are the coefficients of U . UV_0 can be computed from U . Any $m - 1$ shadows cannot solve the system of linear equations and therefore cannot recover the secret.

Advanced Threshold Schemes

The previous examples illustrate only the simplest threshold schemes: Divide a secret into n shadows such that any m can be used to recover the secret. These algorithms can be used to create far more complicated schemes. The following examples will use Shamir's algorithm, although any of the others will work.

To create a scheme in which one person is more important than another, give that person more shadows. If it takes five shadows to recreate a secret and one person has three shadows while everyone else has only one, then that person and two other people can recreate the secret. Without that person, it takes five to recreate the secret.

Two or more people could get multiple shadows. Each person could have a different number of shadows. No matter how the shadows are distributed, any m of them can be used to reconstruct the secret. Someone with $m - 1$ shadows, be it one person or a roomful of people, cannot do it.

In other types of schemes, imagine a scenario with two hostile delegations. You can share the secret so that two people from the 7 in Delegation A and 3 people from the 12 in Delegation B are required to reconstruct the secret. Make a polynomial of degree 3 that is the product of a linear expression and a quadratic expression. Give everyone from Delegation A a shadow that is the result of an evaluation of the linear equation; give everyone from Delegation B a shadow that is the evaluation of the quadratic equation.

Any two shadows from Delegation A can be used to reconstruct the linear equation, but no matter how many other shadows the group has, they cannot get any information about the secret. The same is true for Delegation B: They can get three shadows together to reconstruct the quadratic equation, but they cannot get any

more information necessary to reconstruct the secret. Only when the two delegations share their equations can they be multiplied to reconstruct the secret.

In general, any type of sharing scheme that can be imagined can be implemented. All you have to do is to envision a system of equations that corresponds to the particular scheme. Some excellent papers on generalized secret-sharing schemes are [1462,1463,1464].

Sharing a Secret with Cheaters

This algorithm modifies the standard (m, n) -threshold scheme to detect cheaters [1529]. I demonstrate this using the LaGrange scheme, although it works with the others as well.

Choose a prime, p , that is both larger than n and larger than

$$(s - 1)(m - 1)/e + m$$

where s is the largest possible secret and e is the probability of successful cheating. You can make e as small as you want; it just makes the computation more complex. Construct your shadows as before, except instead of using $1, 2, 3, \dots, n$ for x_i , choose random numbers between 1 and $p - 1$ for x_i .

Now, when Mallory sneaks into the secret reconstruction meeting with his false share, his share has a high probability of not being possible. An impossible secret is, of course, a fake secret. See [1529] for the math.

Unfortunately, while Mallory is exposed as a cheater, he still learns the secret (assuming that there are m other valid shares). Another protocol, from [1529,975], prevents that. The basic idea is to have a series of k secrets, such that none of the participants knows beforehand which is correct. Each secret is larger than the one before, except for the real secret. The participants combine their shadows to generate one secret after the other, until they create a secret that is less than the previous secret. That's the correct one.

This scheme will expose cheaters early, before the secret is generated. There are complications when the participants deliver their shadows one at a time; refer to the papers for details. Other papers on the detection and prevention of cheaters in threshold schemes are [355,114,270].

23.3 SUBLIMINAL CHANNEL

Ong-Schnorr-Shamir

This subliminal channel (see Section 4.2), designed by Gustavus Simmons [1458,1459,1460], uses the Ong-Schnorr-Shamir identification scheme (see Section 20.5). As in the original scheme, the sender (Alice) chooses a public modulus, n , and a private key, k , such that n and k are relatively prime. Unlike the original scheme, k is shared between Alice and Bob, the recipient of the subliminal message.

The public key is calculated:

$$h = -k^2 \bmod n$$

If Alice wants to send the subliminal message M by means of the innocuous message M' , she first confirms that M' and n are relatively prime, and that M and n are relatively prime.

Alice calculates

$$S_1 = 1/2 * ((M'/M + M)) \bmod n$$

$$S_2 = k/2 * ((M'/M - M)) \bmod n$$

Together, the pair, S_1 and S_2 , is the signature under the traditional Ong-Schnorr-Shamir scheme and the carrier of the subliminal message.

Walter the warden (remember him?) can authenticate the message as described by the Ong-Schnorr-Shamir signature scheme, but Bob can do better. He can authenticate the message (it is always possible that Walter can make his own messages). He confirms that

$$S_1^2 - S_2^2/k^2 \equiv M' \pmod{n}$$

If the message is authentic, the receiver can recover the subliminal message using this formula:

$$M = M' / (S_1 + S_2 k^{-1}) \bmod n$$

This works, but remember that the basic Ong-Schnorr-Shamir has been broken.

ElGamal

Simmons's second subliminal channel [1459], described in [1407,1473], is based on the ElGamal signature scheme (see Section 19.6).

Key generation is the same as the basic ElGamal signature scheme. First choose a prime, p , and two random numbers, g and r , such that both g and r are less than p . Then calculate

$$K = g^r \bmod p$$

The public key is K , g , and p . The private key is r . Besides Alice, Bob also knows r ; it is the key that is used to send and read the subliminal message in addition to being the key used to sign the innocuous message.

To send a subliminal message M using the innocuous message M' , M and p must be all relatively prime to each other, and M and $p - 1$ must be relatively prime. Alice calculates

$$X = g^M \bmod p$$

and solves the following equation for Y (using the extended Euclidean algorithm):

$$M' = rX + MY \bmod (p - 1)$$

As in the basic ElGamal scheme, the signature is the pair: X and Y .

Walter can verify the ElGamal signature. He confirms that

$$K^X X^Y \equiv g^{M'} \pmod{p}$$

Bob can recover the subliminal message. First he confirms that

$$(g^r)^X X^Y \equiv g^{M'} \pmod{p}$$

If it does, he accepts the message as genuine (not from Walter).

Then, to recover M , he computes

$$M = (Y^{-1} (M' - rX)) \bmod (p - 1)$$

For example, let $p = 11$ and $g = 2$. The private key, r , is chosen to be 8. This means the public key, which Walter can use to verify the signature, is $g^r \bmod p = 2^8 \bmod 11 = 3$.

To send the subliminal message $M = 9$, using innocuous message $M' = 5$, Alice confirms that 9 and 11 are relatively prime and that 5 and 11 are relatively prime. She also confirms that 9 and $11 - 1 = 10$ are relatively prime. They are, so she calculates

$$X = g^M \bmod p = 2^9 \bmod 11 = 6$$

Then, she solves the following equation for Y :

$$5 = 8 * 6 + 9 * Y \bmod 10$$

$Y = 3$, so the signature is the pair, X and Y : 6 and 3.

Bob confirms that

$$\begin{aligned} (g^r)^X X^Y &\equiv g^{M'} \pmod{p} \\ (2^8)^6 6^3 &\equiv 2^5 \pmod{11} \end{aligned}$$

It does (do the math yourself if you don't trust me), so he then recovers the subliminal message by calculating

$$M = (Y^{-1} (M' - rX)) \bmod (p - 1) = 3^{-1}(5 - 8 * 6) \bmod 10 = 7(7) \bmod 10 = 49 \bmod 10 = 9$$

ESIGN

A subliminal channel can be added to ESIGN [1460] (see Section 20.6).

In ESIGN, the secret key is a pair of large prime numbers, p and q , and the public key is $n = p^2q$. With a subliminal channel, the private key is three primes, p , q , and r , and the public key is n , such that

$$n = p^2qr$$

The variable, r , is the extra piece of information that Bob needs to read the subliminal message.

To sign a normal message, Alice first picks a random number, x , such that x is less than pqr and computes:

$$\begin{aligned} w, \text{ the least integer that is larger than } (H(m) - x^k \bmod n)/pqr \\ s = x + ((w/kx^{k-1}) \bmod p)pqr \end{aligned}$$

$H(m)$ is the hash of the message; k is a security parameter. The value s is the signature.

To verify the signature, Bob computes $s^k \bmod n$. He also computes a , which is the least integer larger than the number of bits of n divided by 3. If $H(m)$ is less than or equal to $s^k \bmod n$, and if $s^k \bmod n$ is less than $H(m) + 2^a$, then the signature is considered valid.

To send a subliminal message, M , using the innocuous message, M' , Alice calcu-

lates s using M in place of $H(m)$. This means that the message must be smaller than p^2qr . She then chooses a random value, u , and calculates

$$x' = M' + ur$$

Then, use this x' value as the “random number” x to sign M' . This second s value is sent as a signature.

Walter can verify that s (the second s) is a valid signature of M' .

Bob can also authenticate the message in the same way. But, since he also knows r , he can calculate

$$s = x' + ypqr = M + ur + ypqr \equiv M \pmod{r}$$

This implementation of a subliminal channel is far better than the previous two. In the Ong-Schnorr-Shamir and ElGamal implementations, Bob has Alice's private key. Besides being able to read subliminal messages from Alice, Bob can impersonate Alice and sign normal documents. Alice can do nothing about this; she must trust Bob to set up this subliminal channel.

The ESIGN scheme doesn't suffer from this problem. Alice's private key is the set of three primes: p , q , and r . Bob's secret key is just r . He knows $n = p^2qr$, but to recover p and q he has to factor that number. If the primes are large enough, Bob has just as much trouble impersonating Alice as would Walter or anyone else.

DSA

There is also a subliminal channel in DSA (see Section 20.1) [1468,1469,1473]. In fact, there are several. The simplest subliminal channel involves the choice of k . It is supposed to be a 160-bit random number. However, if Alice chooses a particular k , then Bob, who also knows Alice's private key, can recover it. Alice can send Bob a 160-bit subliminal message in each DSA signature; everyone else simply verifies Alice's signature. Another complication: Since k should be random, Alice and Bob need to share a one-time pad and encrypt the subliminal message with the one-time pad to generate a k .

DSA has subliminal channels that do not require Bob to know Alice's private key. These also involve choosing particular values of k , but cannot be used to send 160 bits of information. This scheme, presented in [1468,1469], allows Alice and Bob to exchange one bit of subliminal information per signed message.

- (1) Alice and Bob agree on a random prime, P (different from the parameter p in the signature scheme). This is their secret key for the subliminal channel.
- (2) Alice signs an innocuous message, M . If she wants to send Bob the subliminal bit, 1, she makes sure the r parameter of the signature is a quadratic residue modulo P . If she wants to send him a 0, she makes sure the r parameter is a quadratic nonresidue modulo P . She does this by signing the message with random k values until she gets a signature with an r with the requisite property. Since quadratic residues and quadratic nonresidues are equally likely, this shouldn't be too difficult.
- (3) Alice sends the signed message to Bob.

- (4) Bob verifies the signature to make sure the message is authentic. Then he checks whether r is a quadratic residue or a quadratic nonresidue modulo P and recovers the subliminal bit.

Sending multiple bits via this method involves making r either a quadratic residue or a quadratic nonresidue modulo a variety of parameters. See [1468,1469] for details.

This scheme can be easily extended to send multiple subliminal bits per signature. If Alice and Bob agree on two random primes, P and Q , Alice can send two bits by choosing a random k such that r is either a quadratic residue mod P or a quadratic nonresidue mod P , and either a quadratic residue mod Q or a quadratic nonresidue mod Q . A random value of k has a 25 percent chance of producing an r of the correct form.

Here's how Mallory, an unscrupulous implementer of DSA, can have the algorithm leak 10 bits of Alice's private key every time she signs a document.

- (1) Mallory puts his implementation of DSA in a tamperproof VLSI chip, so that no one can examine its inner workings. He creates a 14-bit subliminal channel in his implementation of DSA. That is, he chooses 14 random primes, and has the chip choose a value of k such that r is either a quadratic residue or a quadratic nonresidue modulo each of those 14 primes, depending on the subliminal message.
- (2) Mallory distributes the chips to Alice, Bob, and everyone else who wants them.
- (3) Alice signs a message normally, using her 160-bit private key, x .
- (4) The chip randomly chooses a 10-bit block of x : the first 10 bits, the second 10 bits, and so on. Since there are 16 possible 10-bit blocks, a 4-bit number can identify which block it is. This 4-bit identifier, plus the 10 bits of the key, is the 14-bit subliminal message.
- (5) The chip tries random values of k until it finds one that has the correct quadratic residue properties to send the subliminal message. The odds of a random k being of the correct form are 1 in 16,384. Assuming the chip can test 10,000 values of k per second, it will find one in less than two seconds. This computation does not involve the message and can be performed off-line, before Alice wants to sign a message.
- (6) The chip signs the message normally, using the value of k chosen in step (5).
- (7) Alice sends the digital signature to Bob, or publishes it on the network, or whatever.
- (8) Mallory recovers r and, because he knows the 14 primes, decrypts the subliminal message.

It's scary that even if Alice knows what is happening, she cannot prove it. As long as those 14 secret primes stay secret, Mallory is safe.

Foiling the DSA Subliminal Channel

The subliminal channel relies on the fact that Alice can choose k to transmit subliminal information. To foil the subliminal channel, Alice cannot be allowed to choose k . However, neither can anyone else; if someone else were allowed to choose k , it would allow that person to forge Alice's signature. The only solution is for Alice to jointly generate k with another party, call him Bob, in such a way that Alice cannot control a single bit of k and Bob cannot know a single bit of k . At the end of the protocol, Bob should be able to verify that Alice used the k that they jointly generated.

Here's the protocol [1470,1472,1473]:

- (1) Alice chooses k' and sends Bob

$$u = g^{k'} \bmod p$$

- (2) Bob chooses k'' and sends it to Alice.

- (3) Alice calculates $k = k'k'' \bmod (p-1)$. She uses k to sign her message, M , with the DSA and sends Bob the signature: r and s .

- (4) Bob verifies that

$$((u^{k''} \bmod p) \bmod q) = r$$

If it does, he knows that k was used to sign M .

After step (4), Bob knows that no subliminal information can be embedded in r . If he is a trusted party, he can certify that Alice's signature is subliminal-free. Others will have to trust his certification; Bob cannot prove this fact to a third party with a transcript of the protocol.

A surprising result is that if Bob wants to, he can use this protocol to create his own subliminal channel. Bob can embed a subliminal message in one of Alice's signatures by choosing k'' with certain characteristics. When Simmons discovered this, he dubbed it the "Cuckoo's Channel." Details on how the Cuckoo's Channel works, and a three-pass protocol for generating k that prevents it, are discussed in [1471,1473].

Other Schemes

Any signature scheme can be converted into a subliminal channel [1458,1460,1406]. A protocol for embedding a subliminal channel in the Fiat-Shamir and Feige-Fiat-Shamir protocols, as well as possible abuses of the subliminal channel, can be found in [485].

23.4 UNDENIABLE DIGITAL SIGNATURES

This undeniable signature algorithm (see Section 4.3) is by David Chaum [343,327]. First, a large prime, p , and a primitive element, g , are made public, and used by a group of signers. Alice has a private key, x , and a public key, $g^x \bmod p$.

To sign a message, Alice computes $z = m^x \bmod p$. That's all she has to do. Verification is a little more complicated.

- (1) Bob chooses two random numbers, a and b , both less than p , and sends Alice:

$$c = z^a (g^x)^b \bmod p$$

- (2) Alice computes $t = x^{-1} \bmod (p-1)$, and sends Bob:

$$d = c^t \bmod p$$

- (3) Bob confirms that

$$d \equiv m^a g^b \pmod{p}$$

If it is, he accepts the signature as genuine.

Imagine that Alice and Bob went through this protocol, and Bob is now convinced that Alice signed the message. Bob wants to convince Carol, so he shows her a transcript of the protocol. Dave, however, wants to convince Carol that some other person signed the document. He creates a fake transcript of the protocol. First he generates the message in step (1). Then in step (3) he generates d and the fake transmission from this other person in step (2). Finally, he creates the message in step (2). To Carol, both Bob's and Dave's transcripts are identical. She cannot be convinced of the signature's validity unless she goes through the protocol herself.

Of course, if she were watching over Bob's shoulder as he completed the protocol, she would be convinced. Carol has to see the steps done in order, just as Bob does.

There may be a problem with this signature scheme, but I know of no details. Please pay attention to the literature before you use it.

Another protocol not only has a confirmation protocol—Alice can convince Bob that her signature is valid—but it also has a disavowal protocol; Alice can use a zero-knowledge interactive protocol to convince him that the signature is not valid, if it is not [329].

Like the previous protocol, a group of signers use a shared public large prime, p , and a primitive element, g . Alice has a unique private key, x , and a public key, $g^x \bmod p$. To sign a message, Alice computes $z = m^x \bmod p$.

To verify a signature:

- (1) Bob chooses two random numbers, a and b , both less than p , and sends Alice:

$$c = m^a g^b \bmod p$$

- (2) Alice chooses a random number, q , less than p , and computes and sends to Bob:

$$s_1 = c g^q \bmod p, s_2 = (c g^q)^x \bmod p$$

- (3) Bob sends Alice a and b , so that Alice can confirm that Bob did not cheat in step (1).

- (4) Alice sends Bob q , so that Bob can use m^x and reconstruct s_1 and s_2 . If

$$s_1 \equiv cg^q \pmod{p}$$

$$s_2 \equiv (g^x)^b + qz^a \pmod{p}$$

then the signature is valid.

Alice can also disavow a signature, z , for a message, m . See [329] for details.

Additional protocols for undeniable signatures can be found in [584,344]. Lein Harn and Shoubao Yang proposed a group undeniable signature scheme [700].

Convertible Undeniable Signatures

An algorithm for a **convertible undeniable signature**, which can be verified, disavowed, and also converted to a conventional digital signature is given in [213]. It's based on the ElGamal digital signature algorithm.

Like ElGamal, first choose two primes, p and q , such that q divides $p - 1$. Now you have to create a number, g , less than q . First choose a random number, h , between 2 and $p - 1$. Calculate

$$g = h^{(p-1)/q} \pmod{p}$$

If g equals the 1, choose another random h . If it doesn't, stick with the g you have.

The private keys are two different random numbers, x and z , both less than q . The public keys are p , q , g , y , and u , where

$$y = g^x \pmod{p}$$

$$u = g^z \pmod{p}$$

To compute the convertible undeniable signature of message m (which is actually the hash of a message), first choose a random number, t , between 1 and $q - 1$. Then compute

$$T = g^t \pmod{p}$$

and

$$m' = Ttm \pmod{q}.$$

Now, compute the standard ElGamal signature on m' . Choose a random number, R , such that R is less than and relatively prime to $p - 1$. Then compute $r = g^R \pmod{p}$, and use the extended Euclidean algorithm to compute s , such that

$$m' \equiv rx + Rs \pmod{q}$$

The signature is the ElGamal signature (r, s) , and T .

Here's how Alice verifies her signature to Bob:

- (1) Bob generates two random numbers, a and b . He computes $c = T^{tma}g^b \pmod{p}$ and sends that to Alice.
- (2) Alice generates a random number, k , and computes $h_1 = cg^k \pmod{p}$, and $h_2 = h_1^z \pmod{p}$, and sends both of those numbers to Bob.

- (3) Bob sends Alice a and b .
- (4) Alice verifies that $c = T^{Tma}g^b \bmod p$. She sends k to Bob.
- (5) Bob verifies that $h_1 = T^{Tma}g^{b+k} \bmod p$, and that $h_2 = y^{ta}r^{sa}u^{b+k} \bmod p$.

Alice can convert all of her undeniable signatures to normal signatures by publishing z . Now, anyone can verify her signature without her help.

Undeniable signature schemes can be combined with secret-sharing schemes to create **distributed convertible undeniable signatures** [1235]. Someone can sign a message, then distribute the ability to confirm that the signature is valid. He might, for example, require three out of five people to participate in the protocol in order to convince Bob that the signature is valid. Improvements on this notion deleted the requirement for a trusted dealer [700,1369].

23.5 DESIGNATED CONFIRMER SIGNATURES

Here's how Alice can sign a message and Bob can verify it, such that Carol can verify Alice's signature at some later time to Dave (see Section 4.4) [333].

First, a large prime, p , and a primitive element, g , are made public and used by a group of users. The product of two primes, n , is also public. Carol has a private key, z , and a public key is $h = g^z \bmod p$.

In this protocol Alice can sign m such that Bob is convinced that the signature is valid, but cannot convince a third party.

- (1) Alice chooses a random x and computes

$$a = g^x \bmod p$$

$$b = h^x \bmod p$$

She computes the hash of m , $H(m)$, and the hash of a and b concatenated, $H(a,b)$. She then computes

$$j = (H(m) \oplus H(a, b))^{1/3} \bmod n$$

and sends a , b , and j to Bob.

- (2) Bob chooses two random numbers, s and t , both less than p , and sends Alice

$$c = g^s h^t \bmod p$$

- (3) Alice chooses a random q less than p , and sends Bob

$$d = g^q \bmod p$$

$$e = (cd)^x \bmod p$$

- (4) Bob sends Alice s and t .

- (5) Alice confirms that

$$g^s h^t = c \pmod{p}$$

Then she sends Bob q .

(6) Bob confirms

$$d \equiv g^q \pmod{p}$$

$$e/a^q \equiv a^s b^t \pmod{p}$$

$$H(m) \oplus H(a, b) \equiv j^{1/3} \pmod{n}$$

If they all check out, he accepts the signature as genuine.

Bob cannot use a transcript of this proof to convince Dave that the signature is genuine, but Dave can conduct a protocol with Alice's designated confirmer, Carol. Here's how Carol convinces Dave that a and b constitute a valid signature.

(1) Dave chooses a random u and v , both less than p , and sends Carol

$$k = g^u a^v \pmod{p}$$

(2) Carol chooses a random w , less than p , and sends Dave

$$l = g^w \pmod{p}$$

$$y = \{kl\}^z \pmod{p}$$

(3) Dave sends Carol u and v .

(4) Carol confirms that

$$g^u a^v = k \pmod{p}$$

Then she sends Dave w .

(5) Dave confirms that

$$g^w = l \pmod{p}$$

$$y/h^w = h^u b^v \pmod{p}$$

If they both check out, he accepts the signature as genuine.

In another protocol Carol can convert the designated-confirmer protocol into a conventional digital signature. See [333] for details.

23.6 COMPUTING WITH ENCRYPTED DATA

The Discrete Logarithm Problem

There is a large prime, p , and a generator, g . Alice has a particular value for x , and wants to know e , such that

$$g^e \equiv x \pmod{p}$$

This is a hard problem, and Alice lacks the computational power to compute the result. Bob has the power to solve the problem—he represents the government, or a large computing organization, or whatever. Here's how Bob can do it without Alice revealing x [547,4]:

- (1) Alice chooses a random number, r , less than p .
- (2) Alice computes

$$x' = xg^r \bmod p$$

- (3) Alice asks Bob to solve

$$g^{e'} \equiv x' \pmod{p}$$

- (5) Bob computes e' and sends it to Alice.
- (6) Alice recovers e by computing

$$e = (e' - r) \bmod (p - 1)$$

Similar protocols for the quadratic residuosity problem and for the primitive root problem are in [3,4]. (See also Section 4.8.)

23.7 FAIR COIN FLIPS

The following protocols allow Alice and Bob to flip a fair coin over a data network (see Section 4.9) [194]. This is an example of flipping a coin into a well (see Section 4.10). At first, only Bob knows the result of the coin toss and tells it to Alice. Later, Alice may check to make sure that Bob told her the correct outcome of the toss.

Coin Flipping Using Square Roots

Coin-flip subprotocol:

- (1) Alice chooses two large primes, p and q , and sends their product, n to Bob.
- (2) Bob chooses a random positive integer, r , such that r is less than $n/2$. Bob computes

$$z = r^2 \bmod n$$

and sends z to Alice.

- (3) Alice computes the four square roots of $z \pmod{n}$. She can do this because she knows the factorization of n . Let's call them $+x$, $-x$, $+y$, and $-y$. Call x' the smaller of these two numbers:

$$x \bmod n$$

$$-x \bmod n$$

Similarly, call y' the smaller of these two numbers:

$$y \bmod n$$

$$-y \bmod n$$

Note that r is equal either to x' or y' .

- (4) Alice guesses whether $r = x'$ or $r = y'$, and sends her guess to Bob.

- (5) If Alice's guess is correct, the result of the coin flip is heads. If Alice's guess is incorrect, the result of the coin flip is tails. Bob announces the result of the coin flip.

Verification subprotocol:

- (6) Alice sends p and q to Bob.
 (7) Bob computes x' and y' and sends them to Alice.
 (8) Alice calculates r .

Alice has no way of knowing r , so her guess is real. She only tells Bob one bit of her guess in step (4) to prevent Bob from getting both x' and y' . If Bob has both of those numbers, he can change r after step (4).

Coin Flipping Using Exponentiation Modulo p

Exponentiation modulo a prime number, p , is used as a one-way function in this protocol [1306]:

Coin-flip subprotocol:

- (1) Alice chooses a prime number, p , in such a way that the factorization of $p - 1$ is known and contains at least one large prime.
 (2) Bob selects two primitive elements, h and t , in $\text{GF}(p)$. He sends them to Alice.
 (3) Alice checks that h and t are primitive and then chooses a random integer x , relatively prime to $p - 1$. She then computes one of the two values:

$$y = h^x \bmod p, \text{ or } y = t^x \bmod p$$

She sends y to Bob.

- (4) Bob guesses whether Alice calculated y as a function of h or t , and sends his guess to Alice.
 (5) If Bob's guess is correct, the result of the coin flip is heads. If Bob's guess is incorrect, the result of the coin flip is tails. Alice announces the result of the coin flip.

Verification subprotocol:

- (6) Alice reveals x to Bob. Bob computes $h^x \bmod p$ and $t^x \bmod p$, to confirm that Alice has played fairly and to verify the result of the toss. He also checks that x and $p - 1$ are relatively prime.

For Alice to cheat, she has to know two integers, x and x' , such that $h^x \equiv t^{x'} \pmod{p}$. If she knew those values, she would be able to calculate:

$$\log_t h = x'x^{-1} \bmod p - 1 \text{ and } \log_t h = x^{-1}x' \bmod p - 1$$

These are hard problems.

Alice would be able to do this if she knew $\log_2 h$, but Bob chooses h and t in step (2). Alice has no other recourse except to try to compute the discrete logarithm. Alice could also attempt to cheat by choosing an x that is not relatively prime to $p - 1$, but Bob will detect that in step (6).

Bob can cheat if h and t are not primitive in $\text{GF}(p)$, but Alice can easily check that after step (2) because she knows the prime factorization of $p - 1$.

One nice thing about this protocol is that if Alice and Bob want to flip multiple coins, they can use the same values for p , h , and t . Alice just generates a new x , and the protocol continues from step (3).

Coin Flipping Using Blum Integers

Blum integers can be used in a coin-flipping protocol.

- (1) Alice generates a Blum integer, n , a random x relatively prime to n , $x_0 = x^2 \bmod n$, and $x_1 = x_0^2 \bmod n$. She sends n and x_1 to Bob.
- (2) Bob guesses whether x_0 is even or odd.
- (3) Alice sends x to Bob.
- (4) Bob checks that n is a Blum integer (Alice would have to give Bob the factors of n and proofs of their primality, or execute some zero-knowledge protocol to convince him that n is a Blum integer), and he verifies that $x_0 = x^2 \bmod n$ and $x_1 = x_0^2 \bmod n$. If all this checks out, Bob wins the flip if he guessed correctly.

It is crucial that n be a Blum integer. Otherwise, Alice may be able to find an x'_0 such that $x'^2_0 \bmod n = x^2_0 \bmod n = x_1$, where x'_0 is also a quadratic residue. If x_0 were even and x'_0 were odd (or vice versa), Alice could freely cheat.

23.8 ONE-WAY ACCUMULATORS

There is a simple one-way accumulator function [116] (see Section 4.12):

$$A(x_i, y) = x_{i-1}^y \bmod n$$

The numbers n (n is the product of two primes) and x_0 must be agreed upon in advance. Then, the accumulation of y_1 , y_2 , and y_3 would be

$$((x_0^{y_1} \bmod n)^{y_2} \bmod n)^{y_3} \bmod n$$

This computation is independent of the order of y_1 , y_2 , and y_3 .

23.9 ALL-OR-NOTHING DISCLOSURE OF SECRETS

This protocol allows multiple parties (at least two are required for the protocol to work) to buy individual secrets from a single seller (see Section 4.13) [1374,1175]. First, here's a definition. Take two bit strings, x and y . The **fixed bit index (FBI)** of x and y are the bits where the i th bit of x equals the i th bit of y .

For example:

$$x = 110101001011$$

$$y = 101010000110$$

$$\text{FBI}\{x, y\} = \{1, 4, 5, 11\}$$

(We're reading the bits from right to left, with the right-most bit as zero.)

Now, here's the protocol. Alice is the seller. Bob and Carol are buyers. Alice has k n -bit secrets: S_1, S_2, \dots, S_k . Bob wants to buy secret S_b ; Carol wants to buy secret S_c .

- (1) Alice generates a public-key/private-key key pair and tells Bob (but not Carol) the public key. She generates another public-key/private-key key pair and tells Carol (but not Bob) the public key.
- (2) Bob generates k n -bit random numbers, B_1, B_2, \dots, B_k , and tells them to Carol. Carol generates k n -bit random numbers, C_1, C_2, \dots, C_k , and tells them to Bob.
- (3) Bob encrypts C_b (remember, S_b is the secret he wants to buy) with the public key from Alice. He computes the FBI of C_b and the result he just encrypted. He sends this FBI to Carol.

Carol encrypts B_c (remember, S_c is the secret she wants to buy) with the public key from Alice. She computes the FBI of B_c and the result she just encrypted. She sends this FBI to Bob.

- (4) Bob takes each of the n -bit numbers B_1, B_2, \dots, B_k , and replaces every bit whose index is not in the FBI he received from Carol with its complement. He sends this new list of n -bit numbers, B'_1, B'_2, \dots, B'_k , to Alice.

Carol takes each of the n -bit numbers C_1, C_2, \dots, C_k , and replaces every bit whose index is not in the FBI she received from Bob with its complement. She sends this new list of n -bit numbers, C'_1, C'_2, \dots, C'_k , to Alice.

- (5) Alice decrypts all C'_i with Bob's private key, giving her k n -bit numbers: $C''_1, C''_2, \dots, C''_k$. She computes $S_i \oplus C''_i$ for $i = 1$ to k , and sends the results to Bob.

Alice decrypts all B'_i with Carol's private key, giving her k n -bit numbers: $B''_1, B''_2, \dots, B''_k$. She computes $S_i \oplus B''_i$ for $i = 1$ to k , and sends the results to Carol.

- (6) Bob computes S_b by XORing C_b and the b th number he received from Alice.

Carol computes S_c by XORing B_c and the c th number she received from Alice.

This is complicated. An example will go a long way to help.

Alice has the following eight 12-bit secrets for sale: $S_1 = 1990$, $S_2 = 471$, $S_3 = 3860$, $S_4 = 1487$, $S_5 = 2235$, $S_6 = 3751$, $S_7 = 2546$, and $S_8 = 4043$. Bob wants to buy S_7 . Carol wants to buy S_2 .

- (1) Alice uses the RSA algorithm. The key pair she will use with Bob is: $n = 7387$, $e = 5145$, and $d = 777$. The key pair she will use with Carol is: $n = 2747$, $e = 1421$, and $d = 2261$. She tells Bob and Carol each their public key.
- (2) Bob generates eight 12-bit random numbers, $B_1 = 743$, $B_2 = 1988$, $B_3 = 4001$, $B_4 = 2942$, $B_5 = 3421$, $B_6 = 2210$, $B_7 = 2306$, and $B_8 = 222$, and tells them to Carol. Carol generates eight 12-bit random numbers, $C_1 = 1708$, $C_2 = 711$, $C_3 = 1969$, $C_4 = 3112$, $C_5 = 4014$, $C_6 = 2308$, $C_7 = 2212$, and $C_8 = 222$, and tells them to Bob.
- (3) Bob wants to buy S_7 , so he encrypts C_7 with the public key that Alice gave him.

$$2212^{5145} \bmod 7387 = 5928$$

Now:

$$2212 = 0100010100100$$

$$5928 = 1011100101000$$

So, the FBI of those two numbers is $\{0, 1, 4, 5, 6\}$. He sends this to Carol.

Carol wants to buy S_2 , so she encrypts B_2 with the public key that Alice gave her and computes the FBI of B_2 with the result of her encryption. She sends $\{0, 1, 2, 6, 9, 10\}$ to Bob.

- (4) Bob takes B_1, B_2, \dots, B_8 , and replaces every bit whose index is not in the set $\{0, 1, 2, 6, 9, 10\}$ with its complement. For example:

$$B_2 = 111111000100 = 1988$$

$$B'_2 = 011001111100 = 1660$$

He sends B'_1, B'_2, \dots, B'_8 to Alice.

Carol takes C_1, C_2, \dots, C_8 , and replaces every bit whose index is not in the set $\{0, 1, 4, 5, 6\}$ with its complement. For example:

$$C_7 = 0100010100100 = 2212$$

$$C'_7 = 1011100101000 = 5928$$

She sends C'_1, C'_2, \dots, C'_8 to Alice.

- (5) Alice decrypts all C'_i with Bob's private key and XORs the results with S_i . For example, for $i = 7$:

$$5928^{777} \bmod 7387 = 2212; 2546 \oplus 2212 = 342$$

She sends the results to Bob.

Alice decrypts all B'_i with Carol's private key and XORs the results with S_i . For example, for $i = 2$:

$$1660^{2261} \bmod 2747 = 1988; 471 \oplus 1988 = 1555$$

She sends the results to Carol.

- (6) Bob computes S_7 by XORing C_7 and the seventh number he received from Alice:

$$2212 \oplus 342 = 2546$$

Carol computes S_2 by XORing B_2 and the second number she received from Alice.

$$1988 \oplus 1555 = 471$$

The protocol works for any number of buyers. If Bob, Carol, and Dave want to buy secrets, Alice gives each buyer two public keys, one for each of the others. Each buyer gets a set of numbers from each other buyer. Then, they complete the protocol with Alice for each of their sets of numbers and XOR all of their final results from Alice to get their secret. More details are in [1374,1175].

Unfortunately, a pair of dishonest parties can cheat. Alice and Carol, working together, can easily find out what secret Bob is getting: If they know the FBI of C_b and Bob's encryption algorithm, they can find b such that C_b has the right FBI. And Bob and Carol, working together, can easily get all the secrets from Alice.

If you assume honest parties, there's an easier protocol [389].

- (1) Alice encrypts all of the secrets with RSA and sends them to Bob:

$$C_i = S_i^e \bmod n$$

- (2) Bob chooses his secret C_b , picks a random r , and sends C' to Alice.

$$C' = C_b r^e \bmod n$$

- (3) Alice sends Bob P' .

$$P' = C'^d \bmod n$$

- (4) Bob calculates S_b .

$$S_b = P' r^{-1} \bmod n$$

If the parties may be dishonest, Bob can prove in zero-knowledge that he knows some r such that $C' = C_b r^e \bmod n$ and keep b secret until Alice gives him P' in step (3) [246].

23.10 FAIR AND FAILSAFE CRYPTOSYSTEMS

Fair Diffie-Hellman

Fair cryptosystems are a way to do key escrowing in software (see Section 4.14). This example is from Silvio Micali [1084,1085]. It is patented [1086,1087].

In the basic Diffie-Hellman scheme, a group of users share a prime, p , and a generator, g . Alice's private key is s , and her public key is $t = g^s \bmod p$.

Here's how to make Diffie-Hellman fair (this example uses five trustees).

- (1) Alice chooses five integers, s_1, s_2, s_3, s_4 , and s_5 , each less than $p - 1$. Alice's private key is

$$s = \{s_1 + s_2 + s_3 + s_4 + s_5\} \bmod p - 1$$

and her public key is

$$t = g^s \bmod p$$

Alice also computes

$$t_i = g^{s_i} \bmod p, \text{ for } i = 1 \text{ to } 5$$

Alice's public shares are t_i , and her private shares are s_i .

- (2) Alice sends a private piece and corresponding public piece to each trustee. For example, she sends s_1 and t_1 to trustee 1. She sends t to the KDC.
- (3) Each trustee verifies that

$$t_i = g^{s_i} \bmod p$$

If it does, the trustee signs t_i and sends it to the KDC. The trustee stores s_i in a secure place.

- (4) After receiving all five public pieces, the KDC verifies that

$$t = \{t_1 * t_2 * t_3 * t_4 * t_5\} \bmod p$$

If it does, the KDC approves the public key.

At this point, the KDC knows that the trustees each have a valid piece, and that they can reconstruct the private key if required. However, neither the KDC nor any four of the trustees working together can reconstruct Alice's private key.

Micali's papers [1084,1085] also contain a procedure for making RSA fair and for combining a threshold scheme with the fair cryptosystem, so that m out of n trustees can reconstruct the private key.

Failsafe Diffie-Hellman

Like the previous protocol, a group of users share a prime, p , and a generator, g . Alice's private key is s , and her public key is $t = g^s \bmod p$.

- (1) The KDC chooses a random number, B , between 0 and $p - 2$, and commits to B using a bit-commitment protocol (see Section 4.9).
- (2) Alice chooses a random number, A , between 0 and $p - 2$. She sends $g^A \bmod p$ to the KDC.
- (3) The user "shares" A with each trustee using a verifiable secret-sharing scheme (see Section 3.7).
- (4) The KDC reveals B to Alice.
- (5) Alice verifies the commitment from step (1). Then she sets her public key as

$$t = \{g^A\}g^B \bmod p$$

She sets her private key as

$$s = \{A + B\} \bmod (p - 1)$$

The trustees can reconstruct A . Since the KDC knows B , this is enough to reconstruct s . And Alice cannot make use of any subliminal channels to send unauthorized information. This protocol, discussed in [946,833] is being patented.

23.11 ZERO-KNOWLEDGE PROOFS OF KNOWLEDGE

Zero-Knowledge Proof of a Discrete Logarithm

Peggy wants to prove to Victor that she knows an x that satisfies

$$A^x \equiv B \pmod{p}$$

where p is a prime, and x is a random number relatively prime to $p - 1$. The numbers A , B , and p are public, and x is secret. Here's how Peggy can prove she knows x without revealing it (see Section 5.1) [338,337].

- (1) Peggy generates t random numbers, r_1, r_2, \dots, r_t , where all r_i are less than $p - 1$.
- (2) Peggy computes $h_i = A^{r_i} \pmod{p}$, for all values of i , and sends them to Victor.
- (3) Peggy and Victor engage in a coin-flipping protocol to generate t bits: b_1, b_2, \dots, b_t .
- (4) For all t bits, Peggy does one of the following:
 - a) If $b_i = 0$, she sends Victor r_i
 - b) If $b_i = 1$, she sends Victor $s_i = (r_i - r_j) \pmod{p - 1}$, where j is the lowest value for which $b_j = 1$
- (5) For all t bits, Victor confirms one of the following:
 - a) If $b_i = 0$, that $A^{r_i} \equiv h_i \pmod{p}$
 - b) If $b_i = 1$, that $A^{s_i} \equiv h_i h_j^{-1} \pmod{p}$
- (6) Peggy sends Victor Z , where

$$Z = (x - r_j) \pmod{p - 1}$$
- (7) Victor confirms that

$$A^Z \equiv B h_j^{-1} \pmod{p}$$

Peggy's probability of successfully cheating is 2^{-t} .

Zero-Knowledge Proof of the Ability to Break RSA

Alice knows Carol's private key. Maybe she has broken RSA; maybe she has broken into Carol's house and stolen the key. Alice wants to convince Bob that she knows Carol's key. However, she doesn't want to tell Bob the key or even decrypt one of Carol's messages for Bob. Here's a zero-knowledge protocol by which Alice convinces Bob that she knows Carol's private key [888].

Carol's public key is e , her private key is d , and the RSA modulus is n .

- (1) Alice and Bob agree on a random k and an m such that

$$km \equiv e \pmod{n}$$

They should choose the numbers randomly, using a coin-flip protocol to generate k and then computing m . If both k and m are greater than 3, the protocol continues. Otherwise, they choose again.

- (2) Alice and Bob generate a random ciphertext, C . Again, they should use a coin-flip protocol.
 (3) Alice, using Carol's private key, computes

$$M = C^d \pmod{n}$$

She then computes

$$X = M^k \pmod{n}$$

and sends X to Bob.

- (4) Bob confirms that $X^m \pmod{n} = C$. If it does, he believes Alice.

A similar protocol can be used to demonstrate the ability to break a discrete logarithm problem [888].

Zero-Knowledge Proof that n Is a Blum Integer

There are no known truly practical zero-knowledge proofs that $n = pq$, where p and q are primes congruent to 3 modulo 4. However, if you allow n to be of the form $p'q^s$, where r and s are odd, then the properties which make Blum integers useful in cryptography still hold. And there exists a zero-knowledge proof that n is of that form.

Assume Alice knows the factorization of the Blum integer n , where n is of the form previously discussed. Here's how she can prove to Bob that n is of that form [660].

- (1) Alice sends Bob a number u which has a Jacobi symbol -1 modulo n .
 (2) Alice and Bob jointly agree on random bits: b_1, b_2, \dots, b_k .
 (3) Alice and Bob jointly agree on random numbers: x_1, x_2, \dots, x_k .
 (4) For each $i = 1, 2, \dots, k$, Alice sends Bob a square root modulo n , of one of the four numbers: $x_i, -x_i, ux_i, -ux_i$. The square root must have the Jacobi symbol b_i .

The odds of Alice successfully cheating are one in 2^k .

23.12 BLIND SIGNATURES

The notion of blind signatures (see Section 5.3) was invented by David Chaum [317,323], who also invented their first implementation [318]. It uses the RSA algorithm.

Bob has a public key, e , a private key, d , and a public modulus, n . Alice wants Bob to sign message m blindly.

- (1) Alice chooses a random value, k , between 1 and n . Then she blinds m by computing

$$t = mk^e \bmod n$$

- (2) Bob signs t

$$t^d = (mk^e)^d \bmod n$$

- (3) Alice unblinds t^d by computing

$$s = t^d/k \bmod n$$

- (4) And the result is

$$s = m^d \bmod n$$

This can easily be shown

$$t^d \equiv (mk^e)^d \equiv m^d k \pmod{n}, \text{ so } t^d/k \equiv m^d k/k \equiv m^d \pmod{n}.$$

Chaum invented a family of more complicated blind signature algorithms in [320,324], called blind unanticipated signatures. These signatures are more complex in construction, but more flexible.

23.13 OBLIVIOUS TRANSFER

In this protocol by Michael Rabin [1286], Alice has a 50 percent chance of sending Bob two primes, p , and q . Alice will not know whether the transfer is successful. (See Section 5.5.) (This protocol can be used to send Bob any message with a 50 percent success rate if p and q reveal an RSA private key.)

- (1) Alice sends Bob the product of the two primes: $n = pq$.
- (2) Bob chooses a random x less than n , such that x is relatively prime to n . He sends Alice:

$$a = x^2 \bmod n$$

- (3) Alice, knowing p and q , computes the four roots of a : x , $n - x$, y , and $n - y$. She chooses one of these roots at random and sends it to Bob.
- (4) If Bob receives y or $n - y$, he can compute the greatest common divisor of $x + y$ and n , which is either p or q . Then, of course, $n/p = q$.
If Bob receives x or $n - x$, he can't compute anything.

This protocol may have a weakness: It might be the case that Bob can compute a number a such that given the square root of a you can calculate a factor of n all the time.

23.14 SECURE MULTIPARTY COMPUTATION

This protocol is from [1373]. Alice knows the integer i ; Bob knows the integer j . Alice and Bob together wish to know whether $i \leq j$ or if $i > j$, but neither Alice nor Bob wish to reveal the integer each knows. This special case of secure multiparty computation (see Section 6.2) is sometimes known as **Yao's millionaire problem** [1627].

For this example, assume that i and j range from 1 to 100. Bob has a public key and a private key.

- (1) Alice chooses a large random number, x , and encrypts it in Bob's public key.

$$c = E_B(x)$$

- (2) Alice computes $c - i$ and sends the result to Bob.
- (3) Bob computes the following 100 numbers:

$$y_u = D_B(c - i + u), \text{ for } 1 \leq u \leq 100$$

D_B is the decryption algorithm with Bob's private key.

He chooses a large random prime, p . (The size of p should be somewhat smaller than x . Bob doesn't know x , but Alice could easily tell him the size of x .) He then computes the following 100 numbers:

$$z_u = (y_u \bmod p), \text{ for } 1 \leq u \leq 100$$

He then verifies that, for all $u \neq v$

$$|z_u - z_v| \geq 2$$

and that for all u

$$0 < z_u < p - 1$$

If this is not true, Bob chooses another prime and tries again.

- (4) Bob sends Alice this sequence of numbers in this exact order:

$$z_1, z_2, \dots, z_j, z_{j+1} + 1, z_{j+2} + 1, \dots, z_{100} + 1, p$$

- (5) Alice checks whether the i th number in the sequence is congruent to $x \bmod p$. If it is, she concludes that $i \leq j$. If it is not, she concludes that $i > j$.
- (6) Alice tells Bob the conclusion.

All the verification that Bob goes through in step (3) is to guarantee that no number appears twice in the sequence generated in step (4). Otherwise, if $z_a = z_b$, Alice knows that $a \leq j < b$.

The one drawback to this protocol is that Alice learns the result of the computation before Bob does. Nothing stops her from completing the protocol up to step (5) and then refusing to tell Bob the results in step (6). She could even lie to Bob in step (6).

Example of the Protocol

Assume they're using RSA. Bob's public key is 7 and his private key is 23; $n = 55$. Alice's secret value, i , is 4; Bob's secret value, j , is 2. (Assume that only the values 1, 2, 3, and 4 are possible for i and j .)

- (1) Alice chooses $x = 39$, and $c = E_B(39) = 19$.
- (2) Alice computes $c - i = 19 - 4 = 15$. She sends 15 to Bob.
- (3) Bob computes the following 4 numbers:

$$y_1 = D_B(15 + 1) = 26$$

$$y_2 = D_B(15 + 2) = 18$$

$$y_3 = D_B(15 + 3) = 2$$

$$y_4 = D_B(15 + 4) = 39$$

He chooses $p = 31$ and calculates:

$$z_1 = (26 \bmod 31) = 26$$

$$z_2 = (18 \bmod 31) = 18$$

$$z_3 = (2 \bmod 31) = 2$$

$$z_4 = (39 \bmod 31) = 8$$

He does all the verifications and confirms that the sequence is fine.

- (4) Bob sends Alice this sequence of numbers in this exact order:

$$26, 18, 2 + 1, 8 + 1, 31 = 26, 18, 3, 9, 31$$

- (5) Alice checks whether the 4th number in the sequence is congruent to $x \bmod p$. Since $9 \not\equiv 39 \pmod{31}$, then $i > j$.
- (6) Alice tells Bob.

This protocol can be used to create far more complicated protocols. A group of people can conduct a secret auction over a computer network. They arrange themselves in a logical circle, and through individual pairwise comparisons, determine which is offering the highest price. In order to prevent people from changing their bids in mid-auction, some sort of bit-commitment protocol could also be used. If the auction is a Dutch auction, then the highest bidder gets the item for his highest price. If it is an English auction, then he gets the item for the second-highest price. (This can be determined by a second round of pairwise comparisons.) Similar ideas have applications in bargaining, negotiations, and arbitration.

23.15 PROBABILISTIC ENCRYPTION

The notion of **probabilistic encryption** was invented by Shafi Goldwasser and Silvio Micali [624]. Although its theory makes it the most secure cryptosystem invented, its early implementation was impractical [625]. More recent implementations have changed that.

The point behind probabilistic encryption is to eliminate any information leaked with public-key cryptography. Because a cryptanalyst can always encrypt random messages with a public key, he can get some information. Assuming he has ciphertext $C = E_K(M)$ and is trying to recover plaintext message M , he can pick a random message M' and encrypt it: $C' = E_K(M')$. If $C' = C$, then he guessed the correct plaintext. If it's wrong, he just guesses again.

Also, no partial information is leaked about the original message. With public-key cryptography, sometimes a cryptanalyst can learn things about the bits: The XOR of bits 5, 17, and 39 is 1, and so on. With probabilistic encryption, even this type of information remains hidden.

Not a whole lot of information is to be gained here, but there are potential problems with allowing a cryptanalyst to encrypt random messages with your public key. Some information is being leaked to the cryptanalyst every time he encrypts a message. No one really knows how much.

Probabilistic encryption tries to eliminate that leakage. The goal is that no computation on the ciphertext, or on any other trial plaintexts, can give the cryptanalyst any information about the corresponding plaintext.

In probabilistic encryption, the encrypting algorithm is probabilistic rather than deterministic. In other words, a large number of ciphertexts will decrypt to a given plaintext, and the particular ciphertext used in any given encryption is randomly chosen.

$$C_1 = E_K(M), C_2 = E_K(M), C_3 = E_K(M), \dots, C_i = E_K(M) \\ M = D_K(C_1) = D_K(C_2) = D_K(C_3) = \dots = D_K(C_i)$$

With probabilistic encryption, a cryptanalyst can no longer encrypt random plaintexts looking for the correct ciphertext. To illustrate, assume the cryptanalyst has ciphertext $C_i = E_K(M)$. Even if he guesses M correctly, when he encrypts $E_K(M)$, the result will be a completely different $C: C_j$. He cannot compare C_i and C_j , and so cannot know that he has guessed the message correctly.

This is amazingly cool stuff. Even if a cryptanalyst has the public encryption key, the plaintext, and the ciphertext, he cannot prove that the ciphertext is the encryption of the plaintext without the private decryption key. Even if he tries exhaustive search, he can only prove that every conceivable plaintext is a possible plaintext.

Under this scheme, the ciphertext will always be larger than the plaintext. You can't get around this; it's a result of the fact that many ciphertexts decrypt to the same plaintexts. The first probabilistic encryption scheme [625] resulted in a ciphertext so much larger than the plaintext that it was unusable.

However, Manuel Blum and Goldwasser have an efficient implementation of probabilistic encryption using the Blum Blum Shub (BBS) random-bit generator described in Section 17.9 [199].

The BBS generator is based on the theory of quadratic residues. In English, there are two primes, p and q , that are congruent to 3 modulo 4. That's the private key. Their product, $pq = n$, is the public key. (Mind your ps and qs ; the security of this scheme rests in the difficulty of factoring n .)

To encrypt a message, M , first choose some random x , relatively prime to n . Then compute

$$x_0 = x^2 \bmod n$$

Use x_0 as the seed of the BBS pseudo-random-bit generator and use the output of the generator as a stream cipher. XOR M , one bit at a time, with the output of the generator. The generator spits out bits b_i (the least-significant bit of x_i , where $x_i = x_{i-1}^2 \bmod n$), so

$$M = M_1, M_2, M_3, \dots, M_t$$

$$C = M_1 \oplus b_1, M_2 \oplus b_2, M_3 \oplus b_3, \dots, M_t \oplus b_t$$

where t is the length of the plaintext

Append the last computed value, x_t , to the end of the message and you're done.

The only way to decrypt this message is to recover x_0 and then set up the same BBS generator to XOR with the ciphertext. Because the BBS generator is secure to the left, the value x_t is of no use to the cryptanalyst. Only someone who knows p and q can decrypt the message.

In C, the algorithm to recover x_0 from x_t is:

```
int x0 (int p, int q, int n, int t, int xt)
{
    int a, b, u, v, w, z;

    /* we already know that gcd(p, q) == 1 */
    (void)extended_euclidian(p, q, &a, &b);
    u = modexp ((p+1)/4, t, p-1);
    v = modexp ((q+1)/4, t, q-1);
    w = modexp (xt%p, u, p);
    z = modexp (xt%q, v, q);
    return (b*q*w + a*p*z) % n;
}
```

Once you have x_0 , decryption is easy. Just set up the BBS generator and XOR the output with the ciphertext.

You can make this scheme go even faster by using all the known secure bits of x_i , not just the least significant bit. With this improvement, Blum-Goldwasser probabilistic encryption is faster than RSA while leaking no partial information about the plaintext. You can also prove that the difficulty of breaking this scheme is the same as the difficulty of factoring n .

On the other hand, this scheme is totally insecure against a chosen-ciphertext attack. From the least significant bits of the right quadratic residues, it is possible to calculate the square root of any quadratic residue. If you can do this, then you can factor. For details, consult [1570,1571,35,36].

23.16 QUANTUM CRYPTOGRAPHY

Quantum cryptography taps the natural uncertainty of the quantum world. With it, you can create a communications channel where it is impossible to eavesdrop without disturbing the transmission. The laws of physics secure this quantum channel: even if the eavesdropper can do whatever he wants, even if the eavesdropper has

unlimited computing power, even if $P = NP$. Charles Bennett, Gilles Brassard, Claude Crépeau, and others have expanded on this idea, describing quantum key distribution, quantum coin flipping, quantum bit commitment, quantum oblivious transfer, and quantum secure multiparty computation. Their work is described in [128,129,123,124,125,133,126,394,134,392,243,517,132,130,244,393,396]. The best overview of quantum cryptography can be found in [131]; [1651] is another good nontechnical overview. A complete bibliography of quantum cryptography is [237].

This would still be on the lunatic fringe of cryptography, but Bennett and Brassard actually went and built a working model of the thing [127,121,122]. Now we have *experimental* quantum cryptography.

So sit back, get yourself something to drink, and relax. I'm going to explain what this is all about.

According to quantum mechanics, particles don't actually exist in any single place. They exist in several places at once, with probabilities of being in different places if someone looks. However, it isn't until a scientist comes along and measures the particle that it "collapses" into a single location. But you can't measure every aspect (for example, position and velocity) of a particle at the same time. If you measure one of those two quantities, the very act of measuring it destroys any possibility of measuring the other quantity. The quantum world has a fundamental uncertainty and there's no way to avoid it.

That uncertainty can be used to generate a secret key. As they travel, photons vibrate in some direction; up and down, left to right, or more likely at some angle. Normal sunlight is unpolarized; the photons vibrate every which way. When a large group of photons vibrate in the same direction they are **polarized**. Polarization filters allow only photons that are polarized in a certain direction through; the rest are blocked. For example, a horizontal polarization filter only allows horizontally polarized photons through. Turn that filter 90 degrees, and only vertically polarized photons can come through.

Let's say you have a pulse of horizontally polarized photons. If they try to pass through a horizontally polarized filter, they all get through. Slowly turn that filter 90 degrees; the number of photons getting through gets smaller and smaller, until none get through. This is counterintuitive. You'd think that turning the filter just a little will block all the photons, since the photons are horizontally polarized. But in quantum mechanics, each particle has a probability of suddenly switching its polarization to match the filter. If the angle is a little bit off, it has a high probability. If the angle is 90 degrees off, it has zero probability. And if the angle is 45 degrees off, it has a 50 percent probability of passing through the filter.

Polarization can be measured in any **basis**: two directions at right angles. An example basis is rectilinear: horizontal and vertical. Another is diagonal: left-diagonal and right-diagonal. If a photon pulse is polarized in a given basis and you measure it in the same basis, you learn the polarization. If you measure it in the wrong basis, you get a random result. We're going to use this property to generate a secret key:

- (1) Alice sends Bob a string of photon pulses. Each of the pulses is randomly polarized in one of four directions: horizontal, vertical, left-diagonal, and right-diagonal.

For example, Alice sends Bob:

| / — — \ — | — /

- (2) Bob has a polarization detector. He can set his detector to measure rectilinear polarization or he can set his detector to measure diagonal polarization. He can't do both; quantum mechanics won't let him. Measuring one destroys any possibility of measuring the other. So, he sets his detectors at random, for example:

× + + × × × + × + +

Now, when Bob sets his detector correctly, he will record the correct polarization. If he sets his detector to measure rectilinear polarization and the pulse is polarized rectilinearly, he will learn which way Alice polarized the photon. If he sets his detector to measure diagonal polarization and the pulse is polarized rectilinearly, he will get a random measurement. He won't know the difference. In this example, he might get the result:

/ | — \ / \ — / — |

- (3) Bob tells Alice, over an insecure channel, what settings he used.
 (4) Alice tells Bob which settings were correct. In our example, the detector was correctly set for pulses 2, 6, 7, and 9.
 (5) Alice and Bob keep only those polarizations that were correctly measured. In our example, they keep:

* | * * * \ — * — *

Using a prearranged code, Alice and Bob each translate those polarization measurements into bits. For example, horizontal and left-diagonal might equal one, and vertical and right-diagonal might equal zero. In our example, they both have:

0 0 1 1

So, Alice and Bob have generated four bits. They can generate as many as they like using this system. On the average, Bob will guess the correct setting 50 percent of the time, so Alice has to send $2n$ photon pulses to generate n bits. They can use these bits as a secret key for a symmetric algorithm or they can guarantee perfect secrecy and generate enough bits for a one-time pad.

The really cool thing is that Eve cannot eavesdrop. Just like Bob, she has to guess which type of polarization to measure; and like Bob, half of her guesses will be wrong. Since wrong guesses change the polarization of the photons, she can't help introducing errors in the pulses as she eavesdrops. If she does, Alice and Bob will end up with different bit strings. So, Alice and Bob finish the protocol like this:

- (6) Alice and Bob compare a few bits in their strings. If there are discrepancies, they know they are being bugged. If there are none, they discard the bits they used for comparison and use the rest.

Enhancements to this protocol allow Alice and Bob to use their bits even in the presence of Eve [133,134,192]. They could compare only the parity of subsets of the bits. Then, if no discrepancies are found, they only have to discard one bit of the subset. This detects eavesdropping with only a 50 percent probability, but if they do this with n different subsets Eve's probability of eavesdropping without detection is only 1 in 2^n .

There's no such thing as passive eavesdropping in the quantum world. If Eve tries to recover all the bits, she will necessarily disrupt the communications.

Bennett and Brassard built a working model of quantum key distribution and have exchanged secure bits on a laser table. The latest I heard, some folks at British Telecom were sending bits over a 10-kilometer fiber-optic link [276,1245,1533]. They figure 50 kilometers is feasible. The mind boggles.